

Ravinthiran_Partheepan_Housing_Prediction_using_KNN_Decision_Tree_

October 4, 2023

1 Computational Intelligence and Decision Making - Lab Work 1

Full Name: Ravinthiran Partheepan | **Degree:** Master's in Artificial Intelligence |
Semester: 1st / **Year:** 1st

1.1 KNN Problem Statement:

Problem: based on the given data of historical real estate transactions create the decision-making model (DMM) which aims to predict prices of new real estate objects.

Project workflow:

- P1. Perform given data analysis and preprocessing
- P2. Implement K-Nearest Neighbors (KNN), Decision tree (DT), and random forest (RF) algorithms (You cannot use library functions for these algorithms)
- P3. Use implemented algorithms to create DMM for the given problem and evaluate the results.
- P4. Use “scikit-learn” (or other) library functions for the same algorithms and evaluate the results.
- P5. Write conclusions.

1.2 Library Import

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.neighbors import KNeighborsRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error, \
    mean_absolute_percentage_error
from sklearn.impute import SimpleImputer
```

1.3 Data Import

```
[2]: read_hist_csv = pd.read_csv("C:/Users/ravin/KTU/CI-DM/Task 1/Datasets/Task 1/  
    ↪historicalData.tsv", sep="\t")  
read_hist_csv.head()
```

```
[2]:
```

	Id	LotFrontage	LotArea	Street	Neighborhood	YearBuilt	YearRemodAdd	\
0	1	65.0	8450	Pave	CollgCr	2003	2003	
1	2	80.0	9600	Pave	Veenker	1976	1976	
2	3	68.0	11250	Pave	CollgCr	2001	2002	
3	4	60.0	9550	Pave	Crawfor	1915	1970	
4	5	84.0	14260	Pave	NoRidge	2000	2000	

	CentralAir	PavedDrive	SaleCondition	SalePrice
0	Y	Y	Normal	208500
1	Y	Y	Normal	181500
2	Y	Y	Normal	223500
3	Y	Y	Abnorml	140000
4	Y	Y	Normal	250000

```
[3]: read_new_csv = pd.read_csv("C:/Users/ravin/KTU/CI-DM/Task 1/Datasets/Task 1/  
    ↪newData.tsv", sep="\t")  
read_new_csv.head()
```

```
[3]:
```

	Id	LotFrontage	LotArea	Street	Neighborhood	YearBuilt	YearRemodAdd	\
0	1000	64.0	6762	Pave	CollgCr	2006	2006	
1	1001	74.0	10206	Pave	Edwards	1952	1952	
2	1002	60.0	5400	Pave	OldTown	1920	1950	
3	1003	75.0	11957	Pave	Somerst	2006	2006	
4	1004	NaN	11500	Pave	NWAmes	1976	1976	

	CentralAir	PavedDrive	SaleCondition	SalePrice
0	Y	Y	Normal	206000
1	N	Y	Normal	82000
2	Y	N	Abnorml	86000
3	Y	Y	Normal	232000
4	Y	Y	Normal	136905

1.4 Data type of features: LotFrontage, LotArea, Street,...SalePrice

```
[4]: print(read_hist_csv.dtypes)
```

```
Id                int64  
LotFrontage       float64  
LotArea           int64  
Street            object  
Neighborhood       object  
YearBuilt         int64
```

```
YearRemodAdd      int64
CentralAir         object
PavedDrive         object
SaleCondition      object
SalePrice          int64
dtype: object
```

```
[5]: print(read_new_csv.dtypes)
```

```
Id                int64
LotFrontage       float64
LotArea           int64
Street            object
Neighborhood       object
YearBuilt         int64
YearRemodAdd      int64
CentralAir         object
PavedDrive         object
SaleCondition      object
SalePrice          int64
dtype: object
```

1.5 Data Quality Test

```
[6]: # Checking for Missing values in HistoricalData
check_hist_dat_null = read_hist_csv.isnull().sum()
print(check_hist_dat_null)
```

```
Id                0
LotFrontage       173
LotArea           0
Street            0
Neighborhood       0
YearBuilt         0
YearRemodAdd      0
CentralAir         0
PavedDrive         0
SaleCondition      0
SalePrice          0
dtype: int64
```

```
[7]: # Checking for missing values in NewData
check_new_dat = read_new_csv.isnull().sum()
print(check_new_dat)
```

```
Id                0
LotFrontage       86
LotArea           0
Street            0
```

```

Neighborhood      0
YearBuilt          0
YearRemodAdd      0
CentralAir        0
PavedDrive        0
SaleCondition     0
SalePrice         0
dtype: int64

```

1.6 Missing Value Imputation

```

[8]: # Historical Data NaN imputation
# Dropping the rows with NaN would eliminate large portions of data because
# ↪ since there are 173 cells in the LotFrontage features has NaN
# So imputing those missing values with value 0

fill_na_hist = read_hist_csv.fillna(read_hist_csv.median()) #Median
check_Hist_NaN = fill_na_hist.isnull().sum()
print(check_Hist_NaN)

```

```

Id              0
LotFrontage     0
LotArea         0
Street          0
Neighborhood    0
YearBuilt       0
YearRemodAdd    0
CentralAir      0
PavedDrive      0
SaleCondition   0
SalePrice       0
dtype: int64

```

C:\Users\ravin\AppData\Local\Temp\ipykernel_11844\2933457410.py:5:
FutureWarning: The default value of numeric_only in DataFrame.median is deprecated. In a future version, it will default to False. In addition, specifying 'numeric_only=None' is deprecated. Select only valid columns or specify the value of numeric_only to silence this warning.

```

fill_na_hist = read_hist_csv.fillna(read_hist_csv.median()) #Median

```

```

[9]: # NewData NaN Imputation
fill_na_new = read_new_csv.fillna(0)
check_New_NaN = fill_na_new.isnull().sum()
print(check_New_NaN)

```

```

Id              0
LotFrontage     0
LotArea         0
Street          0

```

```

Neighborhood    0
YearBuilt       0
YearRemodAdd    0
CentralAir      0
PavedDrive      0
SaleCondition   0
SalePrice       0
dtype: int64

```

```

[10]: # Dropping the ID field
read_hist_csv.drop(['Id'], axis=1, inplace=True)

```

1.7 Correlation Analysis

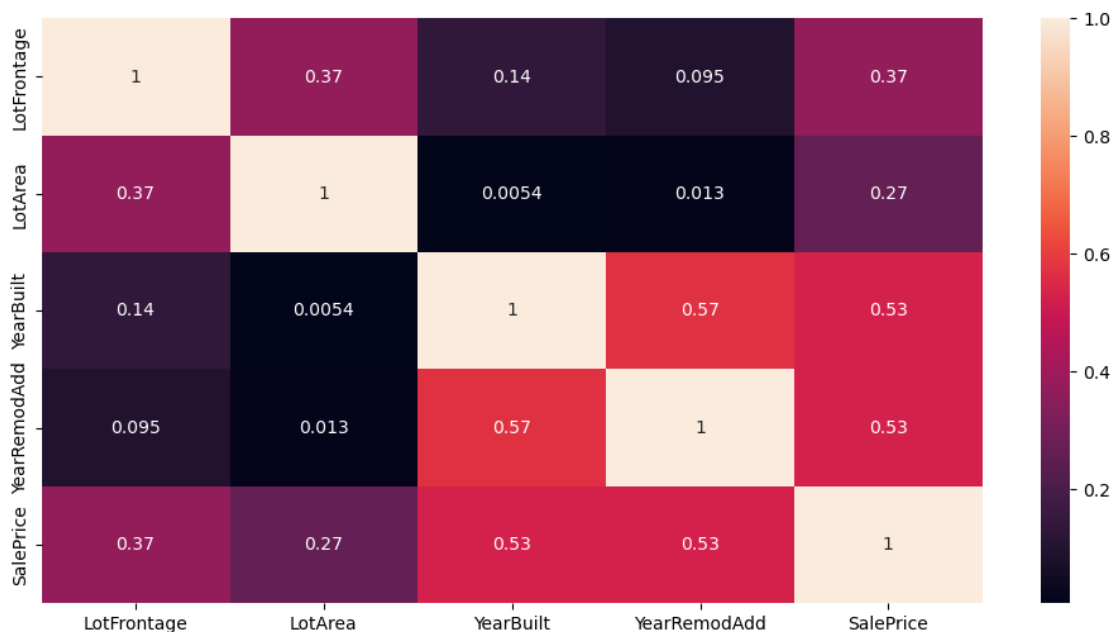
```

[11]: # Implementing correlation to understand the dependency between each feature,
      ↪ whether it is positively correlated or negative correlated
plt.figure(figsize=(12,6))
sns.heatmap(read_hist_csv.corr(), annot=True)

```

C:\Users\ravin\AppData\Local\Temp\ipykernel_11844\1103431315.py:3:
FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.
sns.heatmap(read_hist_csv.corr(), annot=True)

[11]: <Axes: >



1.8 Checking for Duplicates

```
[12]: # Histroical Data Feature Duplication Check
duplicate_hist = read_hist_csv[read_hist_csv.duplicated()]
print(duplicate_hist)
```

Empty DataFrame

Columns: [LotFrontage, LotArea, Street, Neighborhood, YearBuilt, YearRemodAdd, CentralAir, PavedDrive, SaleCondition, SalePrice]

Index: []

```
[13]: duplicate_new = read_new_csv[read_new_csv.duplicated()]
print(duplicate_new)
```

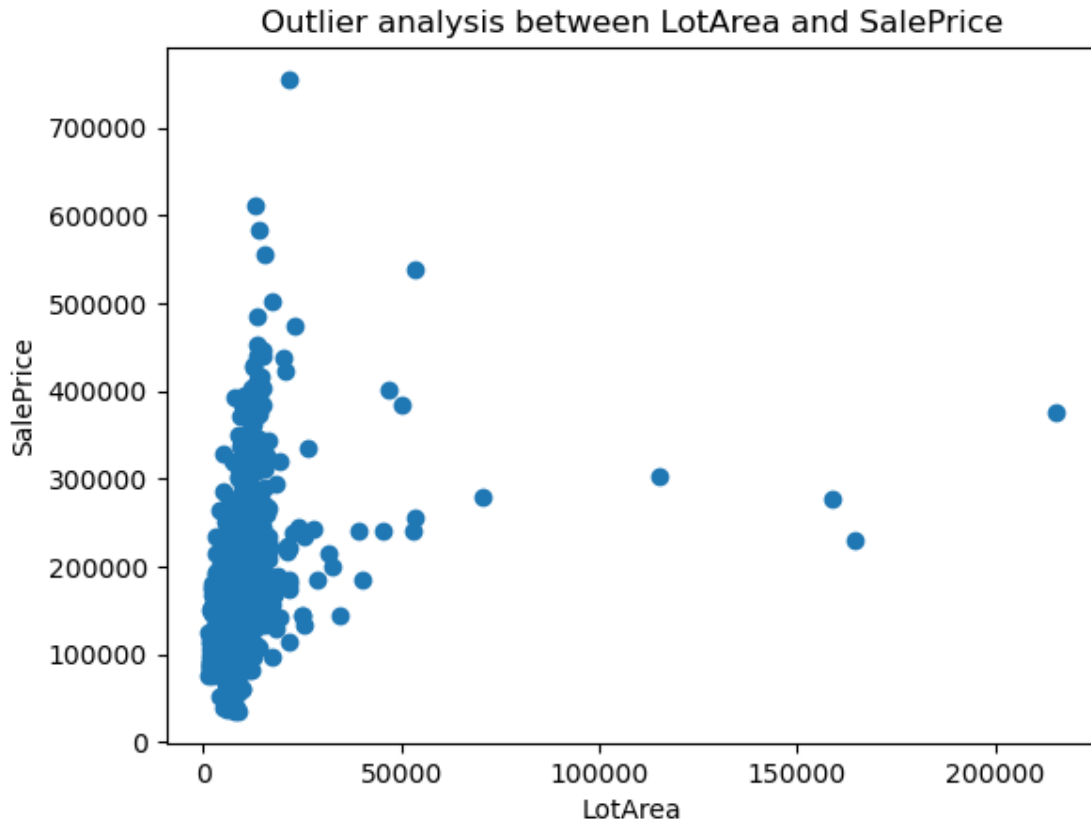
Empty DataFrame

Columns: [Id, LotFrontage, LotArea, Street, Neighborhood, YearBuilt, YearRemodAdd, CentralAir, PavedDrive, SaleCondition, SalePrice]

Index: []

1.9 Outlier Analysis

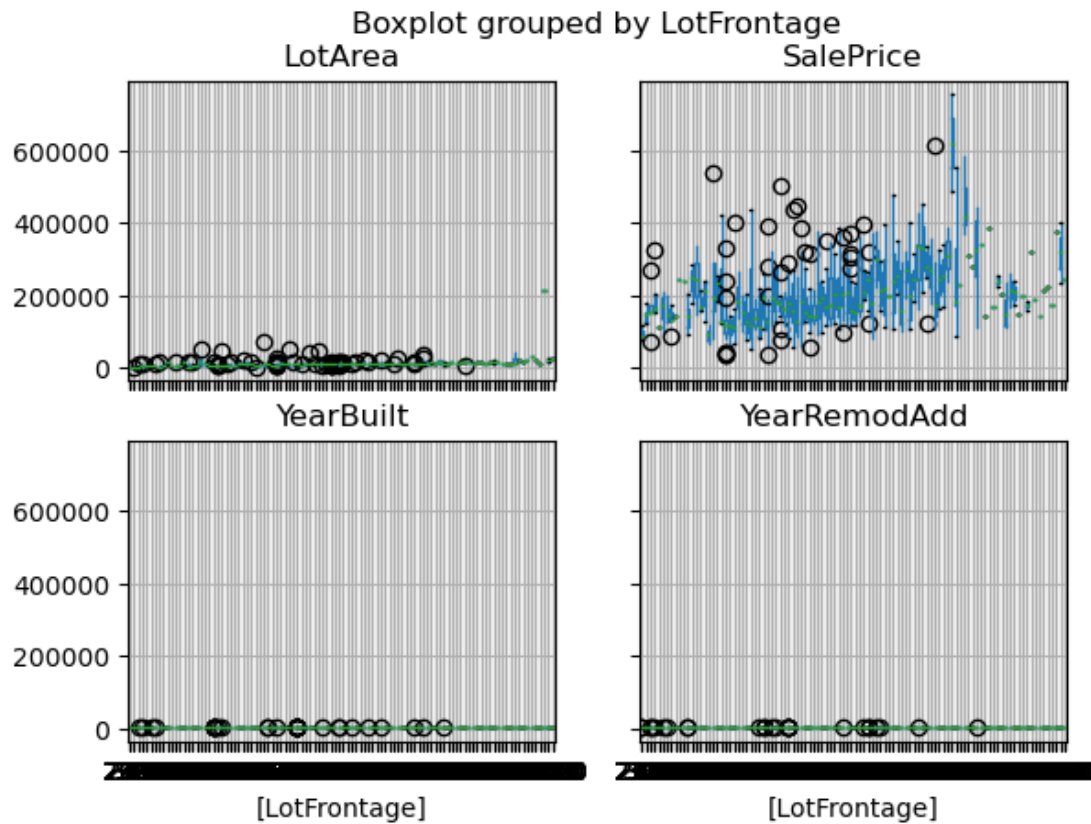
```
[14]: scat_outlier = plt.scatter(read_hist_csv["LotArea"], read_hist_csv["SalePrice"])
plt.title("Outlier analysis between LotArea and SalePrice")
plt.xlabel("LotArea")
plt.ylabel("SalePrice")
plt.show()
```



```
[15]: # Outlier analysis using Box-Plot
scale_fig = plt.figure(figsize=(10,7))
read_hist_csv.boxplot(by="LotFrontage")
```

```
[15]: array([[<Axes: title={'center': 'LotArea'}, xlabel=' [LotFrontage] '>,
    <Axes: title={'center': 'SalePrice'}, xlabel=' [LotFrontage] '>,
    [<Axes: title={'center': 'YearBuilt'}, xlabel=' [LotFrontage] '>,
    <Axes: title={'center': 'YearRemodAdd'}, xlabel=' [LotFrontage] '>]],
    dtype=object)
```

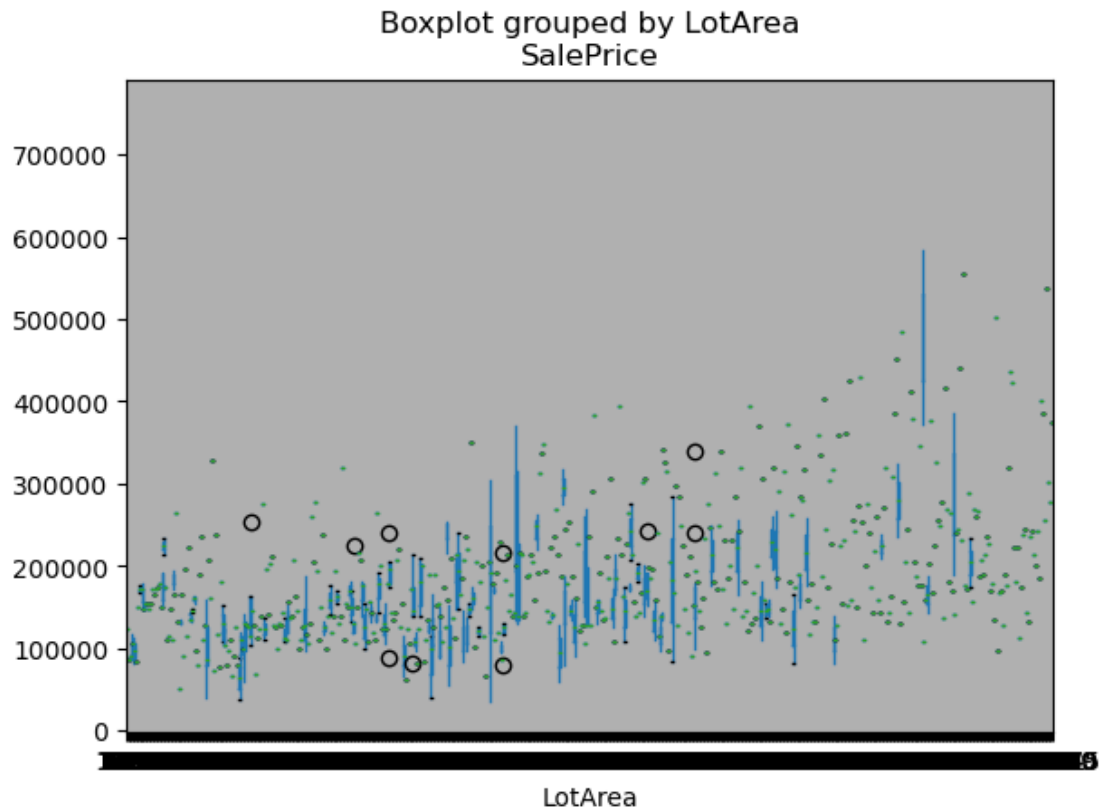
<Figure size 1000x700 with 0 Axes>



```
[16]: scale_fig = plt.figure(figsize=(10,7))
      read_hist_csv.boxplot(by="LotArea", widths= 1.0, column=["SalePrice"])
```

```
[16]: <Axes: title={'center': 'SalePrice'}, xlabel='LotArea'>
```

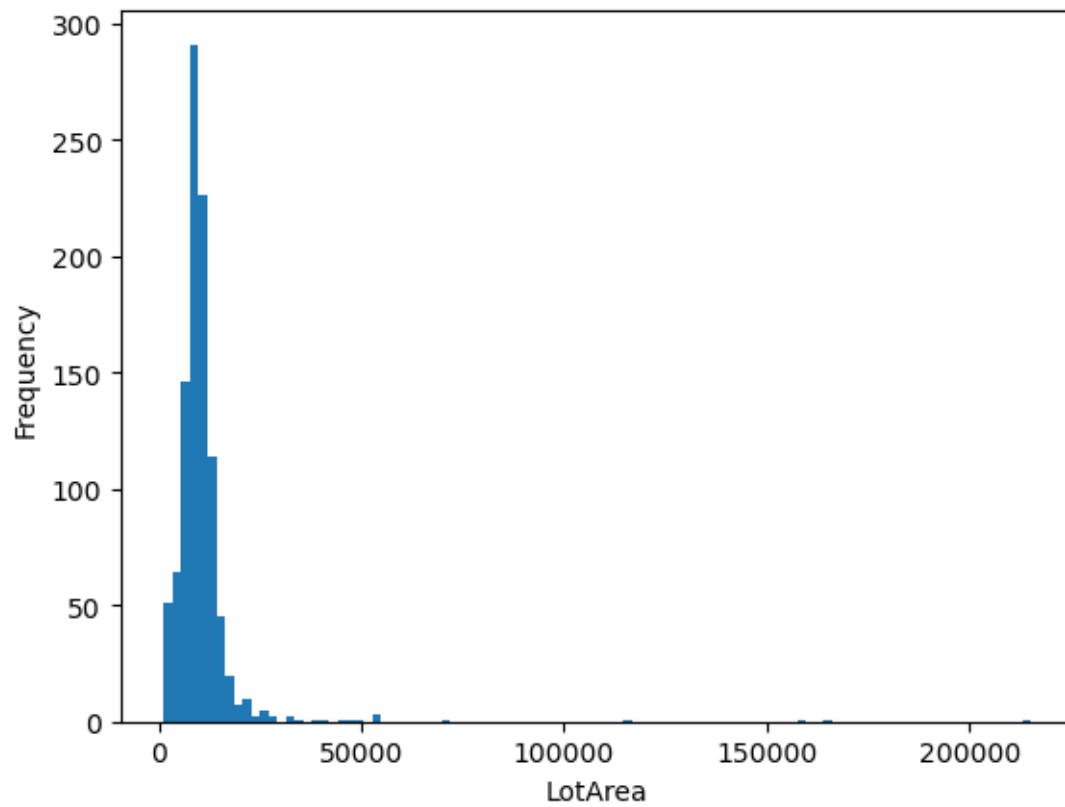
```
<Figure size 1000x700 with 0 Axes>
```

1.10 Data Distribution

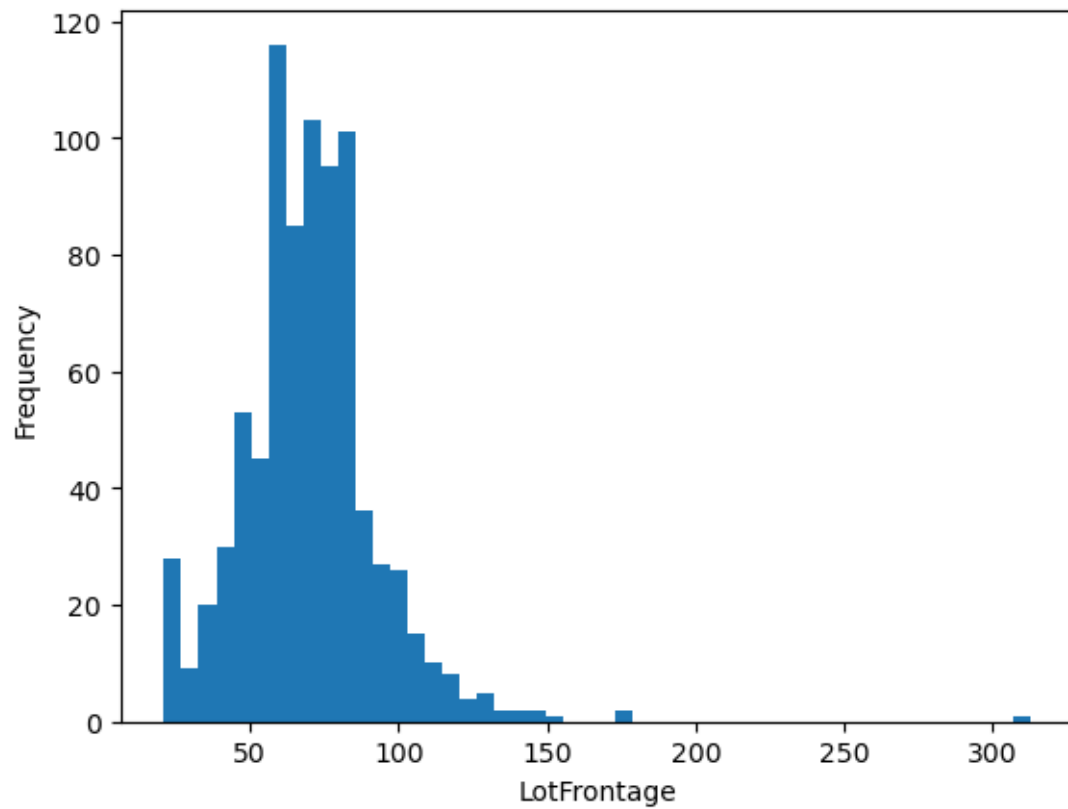
```
[17]: plt.hist(read_hist_csv["LotArea"], bins=100)
      plt.xlabel("LotArea")
      plt.ylabel("Frequency")
```

```
[17]: Text(0, 0.5, 'Frequency')
```



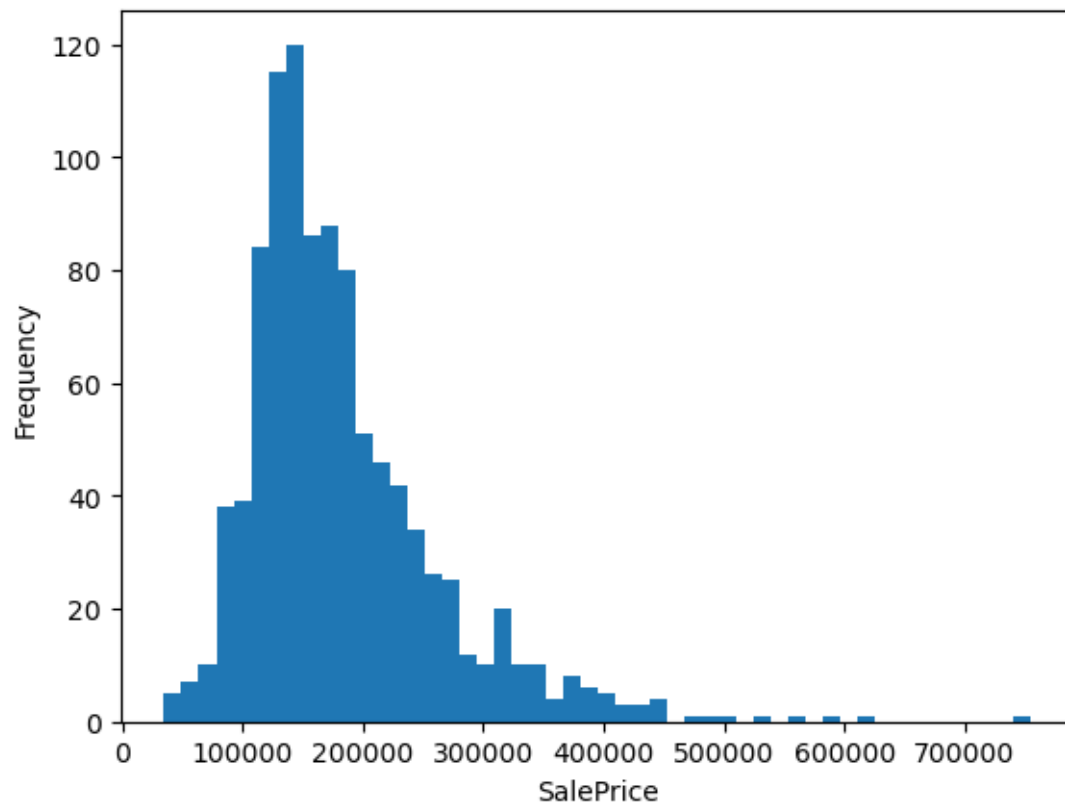
```
[18]: plt.hist(read_hist_csv["LotFrontage"], bins=50)
plt.xlabel("LotFrontage")
plt.ylabel("Frequency")
```

```
[18]: Text(0, 0.5, 'Frequency')
```



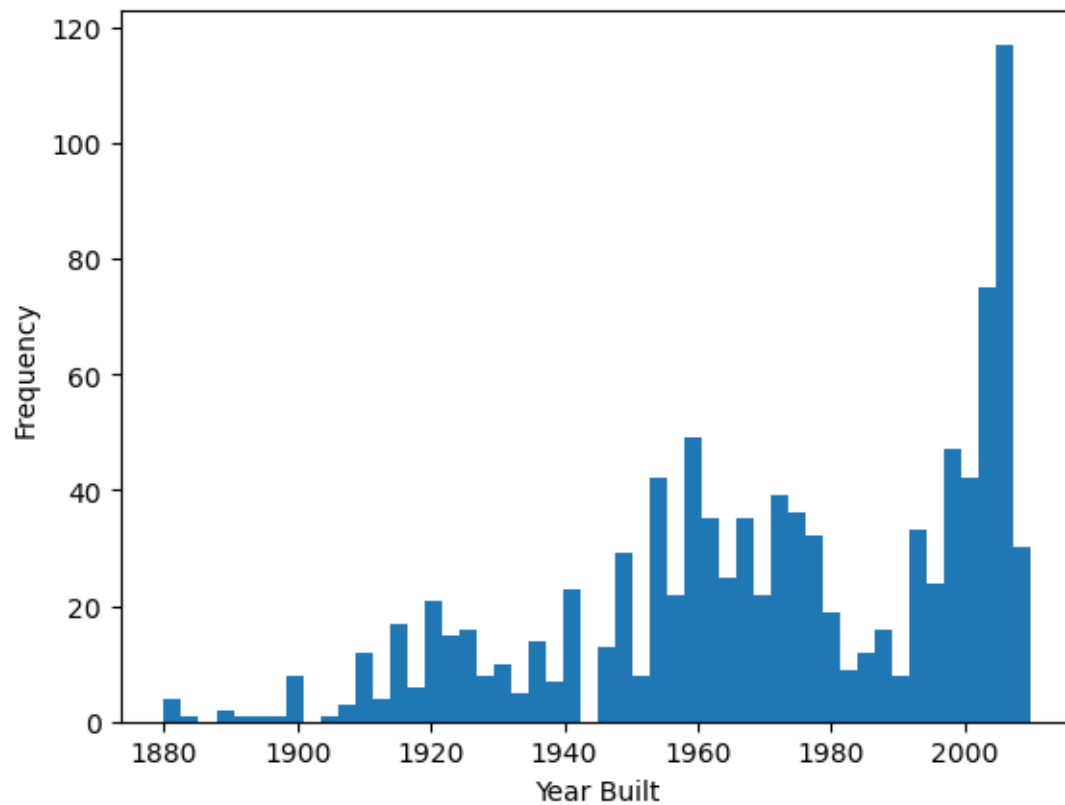
```
[19]: plt.hist(read_hist_csv["SalePrice"], bins=50)
      plt.xlabel("SalePrice")
      plt.ylabel("Frequency")
```

```
[19]: Text(0, 0.5, 'Frequency')
```



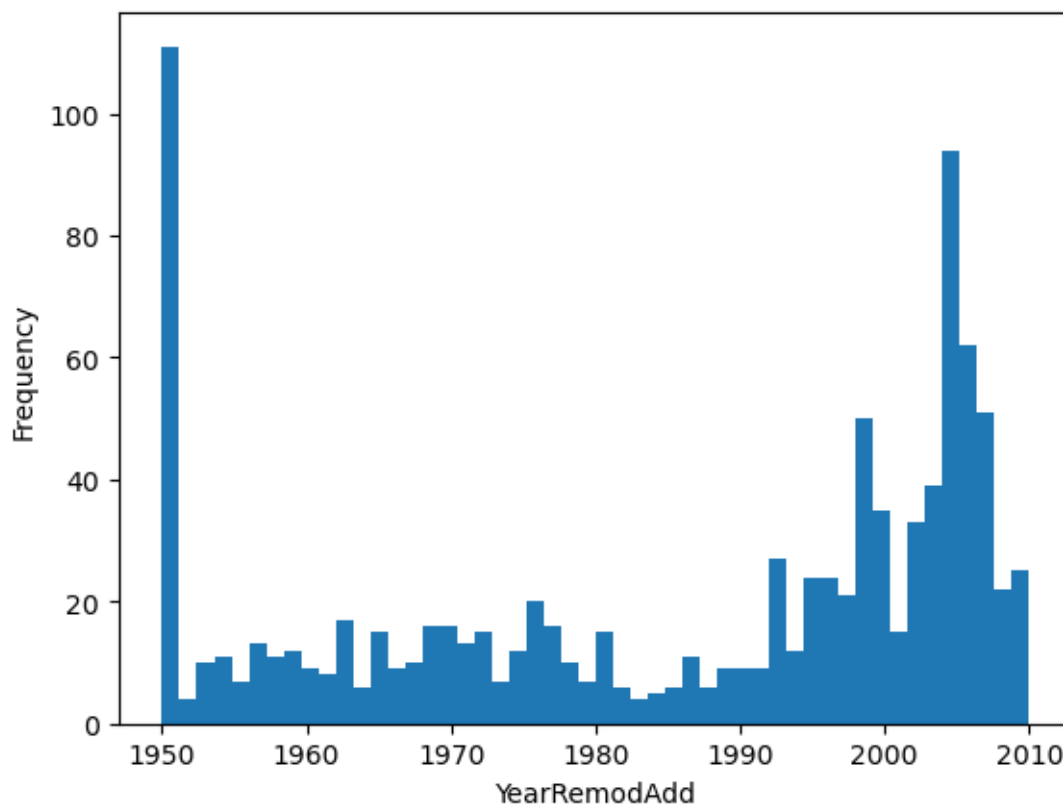
```
[20]: plt.hist(read_hist_csv["YearBuilt"], bins=50)
plt.xlabel("Year Built")
plt.ylabel("Frequency")
```

```
[20]: Text(0, 0.5, 'Frequency')
```



```
[21]: plt.hist(read_hist_csv["YearRemodAdd"], bins=50)
plt.xlabel("YearRemodAdd")
plt.ylabel("Frequency")
```

```
[21]: Text(0, 0.5, 'Frequency')
```



1.10.1 Derived Feature between [YearBuilt - YearRemodAdd]

```
[22]: # The features = [YearBuilt, YearRemoteAdd] share strong positive corrleation,
      ↪with value 0.57
      # We can create derived features for finding the difference between the two
      ↪features using [YearBuilt-YearRemoteAdd]
      # Note: RemoteYearAdd feature states the year of rennovation / remodelling
      ↪performed

      read_hist_csv["derived_features"] = read_hist_csv["YearRemodAdd"] -
      ↪read_hist_csv["YearBuilt"]
      print(read_hist_csv["derived_features"])
```

```
0      0
1      0
2      1
3     55
4      0
...
994    1
995    4
```

```

996      0
997      0
998     28
Name: derived_features, Length: 999, dtype: int64

```

```

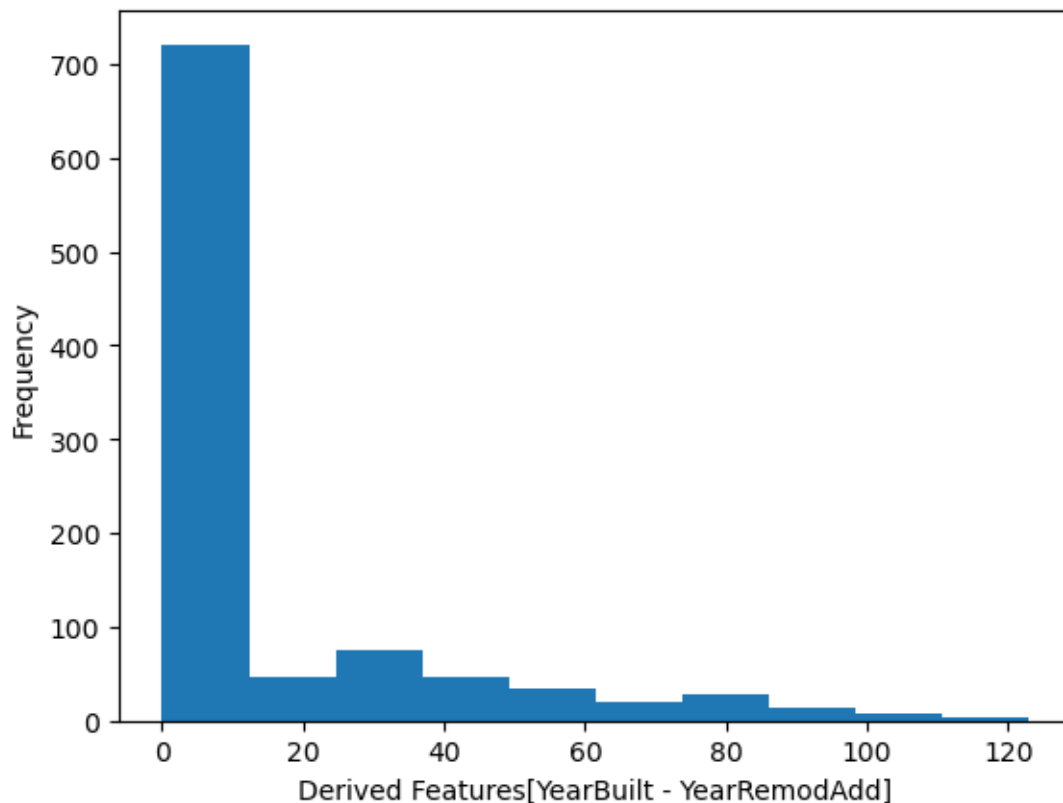
[23]: plt.hist(read_hist_csv["derived_features"], bins=10)
      plt.xlabel("Derived Features[YearBuilt - YearRemodAdd]")
      plt.ylabel("Frequency")

```

```

[23]: Text(0, 0.5, 'Frequency')

```



1.11 Train-Test Dataset Split

```

[24]: features = ['LotFrontage', 'LotArea', 'Street', 'Neighborhood', 'YearBuilt', '
      ↪ 'YearRemodAdd', 'CentralAir', 'PavedDrive', 'SaleCondition']
      X = read_hist_csv[features]
      y = read_hist_csv['SalePrice']

```

```

[25]: # Missing values
      feature_na = ['LotFrontage']
      numeric_imputer = SimpleImputer(strategy='median')

```

```

for elements in feature_na:
    X[elements] = numeric_imputer.fit_transform(X[[elements]])

check_new_dat = X.isnull().sum()
print(check_new_dat)

```

```

LotFrontage      0
LotArea          0
Street           0
Neighborhood     0
YearBuilt        0
YearRemodAdd     0
CentralAir       0
PavedDrive       0
SaleCondition    0
dtype: int64

```

C:\Users\ravin\AppData\Local\Temp\ipykernel_11844\2595087311.py:5:
 SettingWithCopyWarning:
 A value is trying to be set on a copy of a slice from a DataFrame.
 Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
 X[elements] = numeric_imputer.fit_transform(X[[elements]])

1.11.1 Encoding categorical features

```

[26]: label_encoders = {}
categorical_columns = ['Street', 'Neighborhood', 'CentralAir', 'PavedDrive', 'SaleCondition']
for col in categorical_columns:
    lab_encode = LabelEncoder()
    X[col] = lab_encode.fit_transform(X[col])
    label_encoders[col] = lab_encode

```

C:\Users\ravin\AppData\Local\Temp\ipykernel_11844\377571960.py:5:
 SettingWithCopyWarning:
 A value is trying to be set on a copy of a slice from a DataFrame.
 Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
 X[col] = lab_encode.fit_transform(X[col])
 C:\Users\ravin\AppData\Local\Temp\ipykernel_11844\377571960.py:5:
 SettingWithCopyWarning:
 A value is trying to be set on a copy of a slice from a DataFrame.
 Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
X[col] = lab_encode.fit_transform(X[col])
```

C:\Users\ravin\AppData\Local\Temp\ipykernel_11844\377571960.py:5:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
X[col] = lab_encode.fit_transform(X[col])
```

C:\Users\ravin\AppData\Local\Temp\ipykernel_11844\377571960.py:5:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
X[col] = lab_encode.fit_transform(X[col])
```

C:\Users\ravin\AppData\Local\Temp\ipykernel_11844\377571960.py:5:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
X[col] = lab_encode.fit_transform(X[col])
```

```
[27]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
```

1.12 Normalization for Features = [LotArea, LotFrontage, SalesPrice,..]

```
[28]: # Using standard scaler to normalize kurtosis - Sales price
# LotArea and LotFrontage Follows Gaussian Distribution so there's no need to
↳ perform normalization
# Since the Price feature follow kurtosis we use log-scaling

normalize = StandardScaler() # Z-Score Normalization
X_train = normalize.fit_transform(X_train)
X_test = normalize.transform(X_test)
```

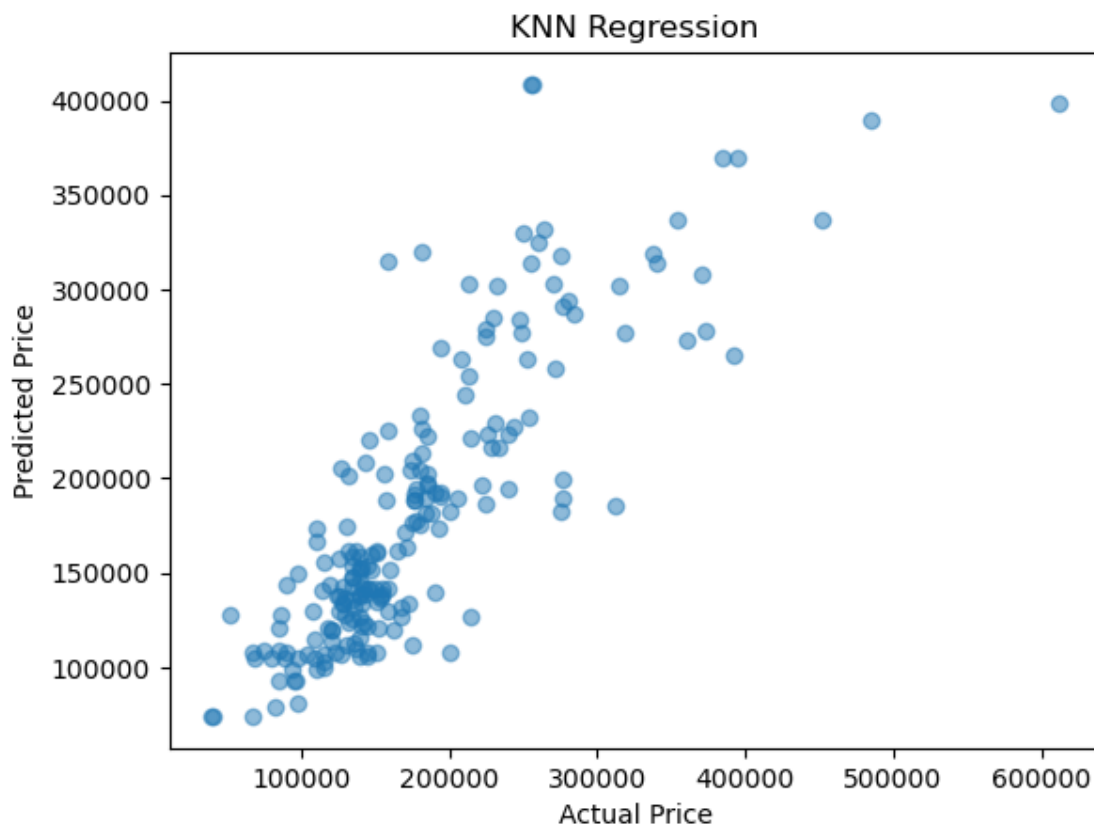
1.13 P2. Implementation of KNN

```
[29]: knn_regressor = KNeighborsRegressor(n_neighbors=5)
      knn_regressor.fit(X_train, y_train)

      # Making predictions on the test dataset
      knn_predictions = knn_regressor.predict(X_test)
```

```
[30]: plt.subplot()
      plt.scatter(y_test, knn_predictions, alpha=0.5)
      plt.xlabel("Actual Price")
      plt.ylabel("Predicted Price")
      plt.title("KNN Regression")
```

```
[30]: Text(0.5, 1.0, 'KNN Regression')
```

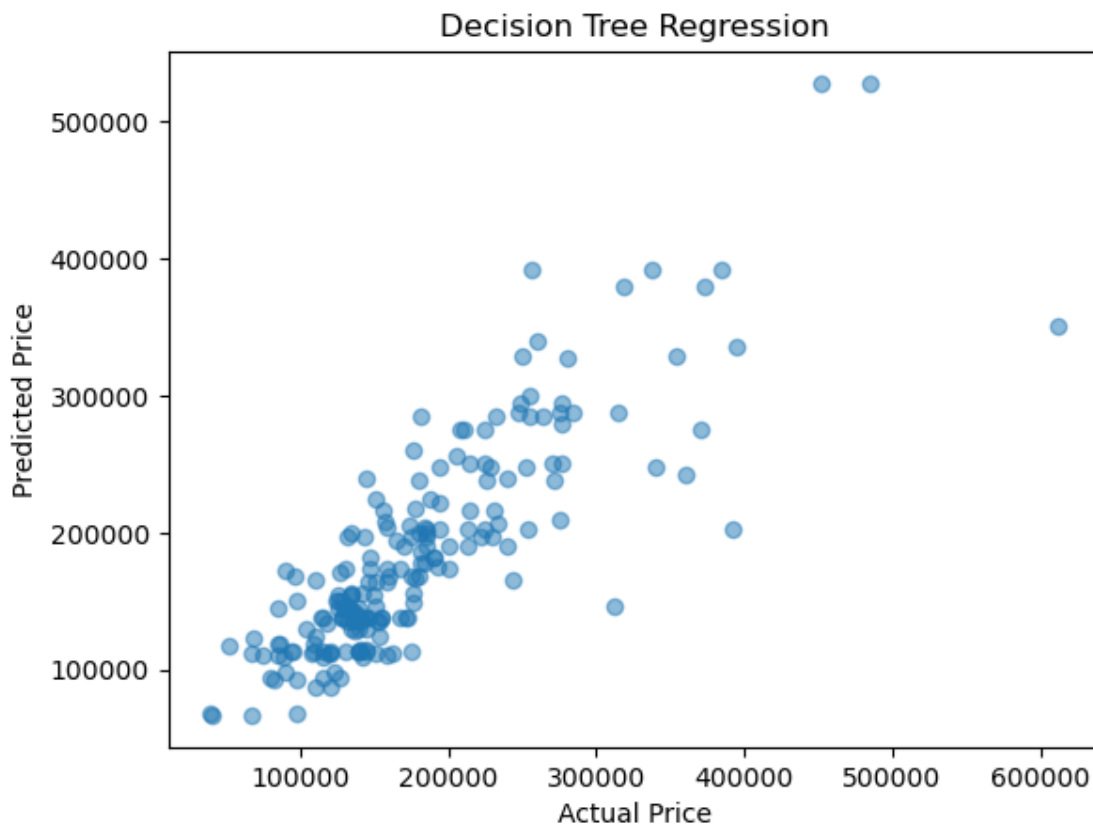


1.14 P2. Implementation of Decision Tree

```
[31]: decision_tree_regressor = DecisionTreeRegressor(min_samples_leaf=3,  
    ↪max_depth=10, random_state=42)  
decision_tree_regressor.fit(X_train, y_train)  
  
# Making predictions on the test set  
dt_predictions = decision_tree_regressor.predict(X_test)
```

```
[32]: plt.subplot()  
plt.scatter(y_test, dt_predictions, alpha=0.5)  
plt.xlabel("Actual Price")  
plt.ylabel("Predicted Price")  
plt.title("Decision Tree Regression")
```

```
[32]: Text(0.5, 1.0, 'Decision Tree Regression')
```

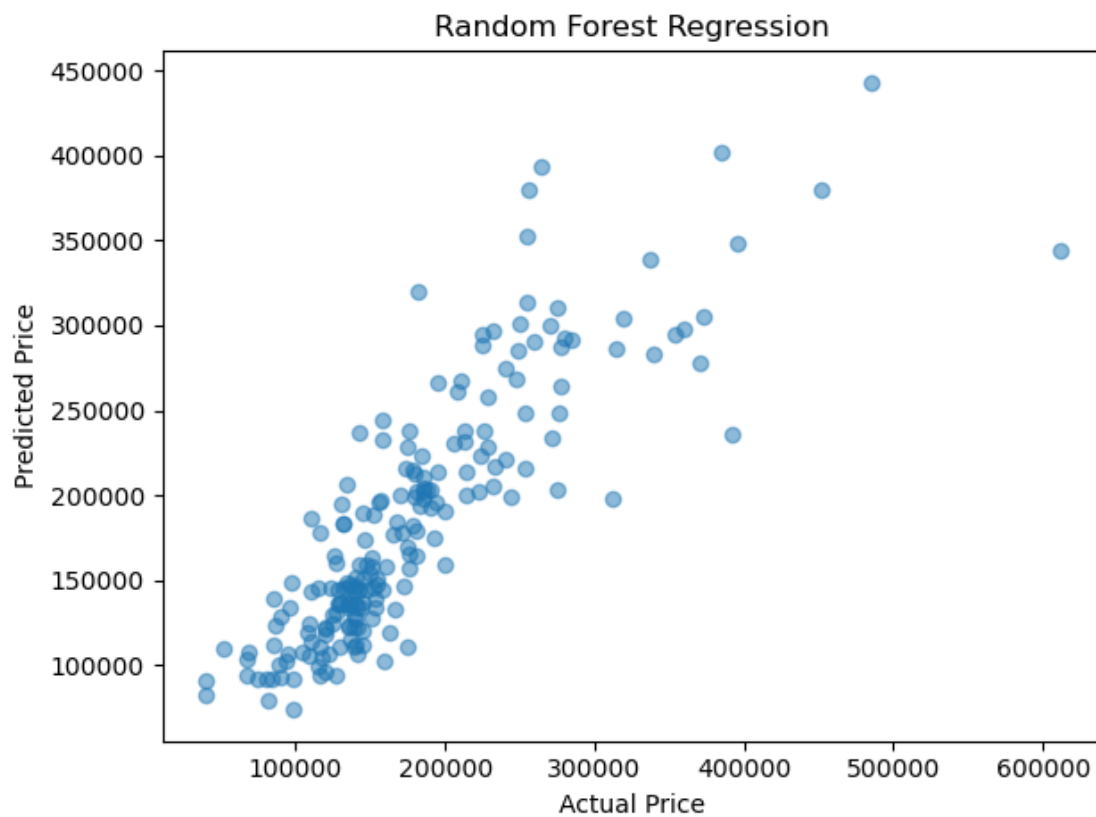


1.15 P2. Implementation of Random Forest

```
[33]: random_forest_regressor = RandomForestRegressor(n_estimators=100, random_state=42)
      random_forest_regressor.fit(X_train, y_train)

      # Making predictions on the test set
      rf_predictions = random_forest_regressor.predict(X_test)
```

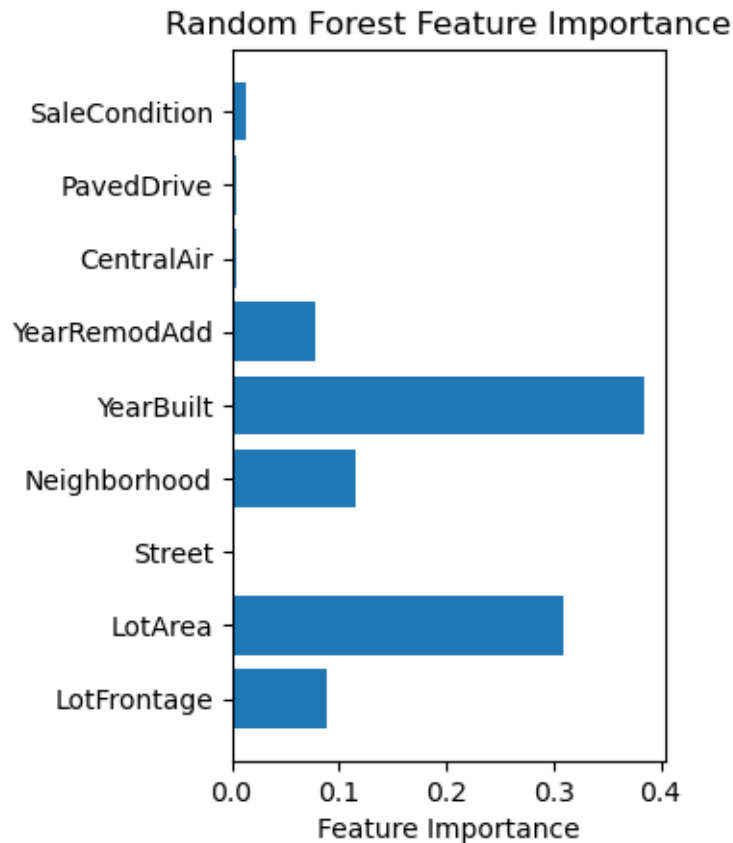
```
[34]: plt.subplot()
      plt.scatter(y_test, rf_predictions, alpha=0.5)
      plt.xlabel("Actual Price")
      plt.ylabel("Predicted Price")
      plt.title("Random Forest Regression")
      plt.tight_layout()
      plt.show()
```



1.16 Feature Importance Analysis

```
[35]: rf_feature_importance = random_forest_regressor.feature_importances_  
plt.subplot(122)  
plt.barh(features, rf_feature_importance)  
plt.xlabel("Feature Importance")  
plt.title("Random Forest Feature Importance")
```

```
[35]: Text(0.5, 1.0, 'Random Forest Feature Importance')
```



1.17 Evaluation models using MAE, MAPE metrics

```
[37]: # MAE and MAPE evaluation of KNN Model  
knn_mae = mean_absolute_error(y_test, knn_predictions)  
knn_mape = mean_absolute_percentage_error(y_test, knn_predictions)  
  
print("KNN Regressor MAE: ", knn_mae)  
print("KNN Regressor MAPE", knn_mape)  
  
# MAE and MAPE evaluation of Decision Tree Model
```

```

dt_mae = mean_absolute_error(y_test, dt_predictions)
dt_mape = mean_absolute_percentage_error(y_test, dt_predictions)

print("DT Regressor MAE: ", dt_mae)
print("DTRegressor MAPE", dt_mape)

# MAE and MAPE evaluation of Random Forest Model
rf_mae = mean_absolute_error(y_test, rf_predictions)
rf_mape = mean_absolute_percentage_error(y_test, rf_predictions)

print("RF Regressor MAE: ", rf_mae)
print("RF Regressor MAPE", rf_mape)

# Mean Absolute Error Visualization
trained_models = ["KNN", "Decision Tree", "Random Forest"]
mae_scores = [knn_mae, dt_mae, rf_mae]

plt.subplot()
plt.bar(trained_models, mae_scores)
plt.ylabel("MAE")
plt.title("Mean Absolute Error")

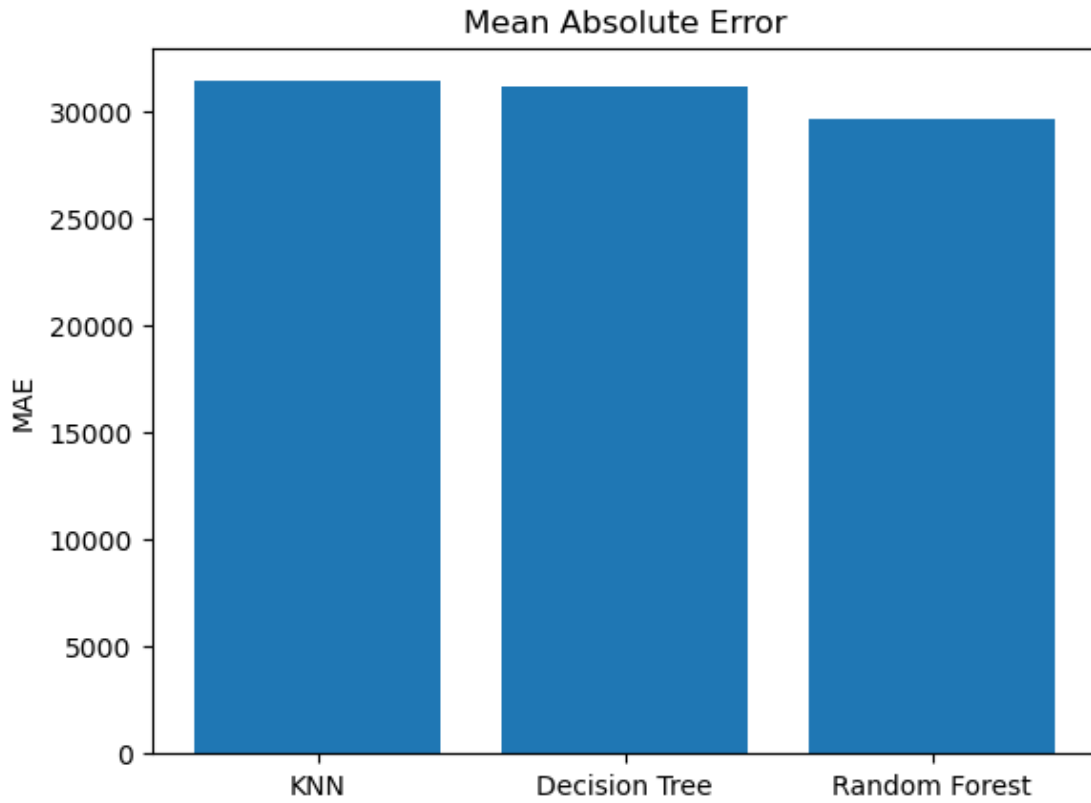
```

```

KNN Regressor MAE: 31419.286
KNN Regressor MAPE 0.1869290327180879
DT Regressor MAE: 31206.244614645097
DTRegressor MAPE 0.1899427934301395
RF Regressor MAE: 29650.54651095238
RF Regressor MAPE 0.1815675752904022

```

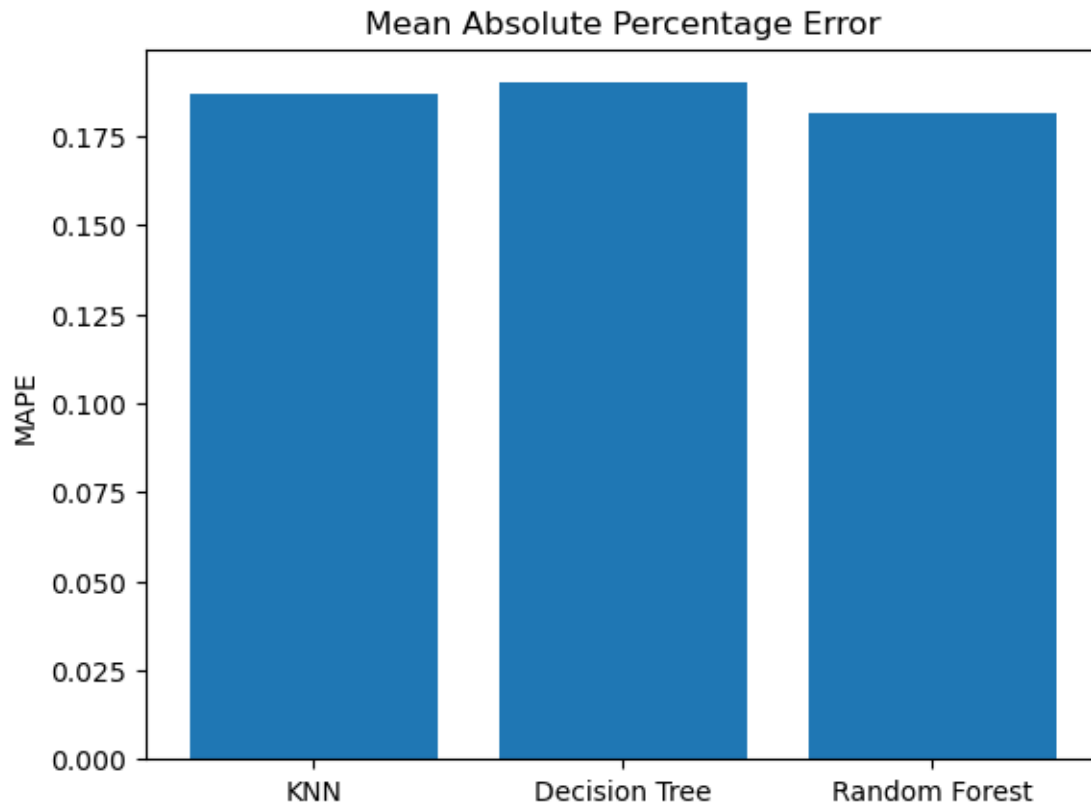
```
[37]: Text(0.5, 1.0, 'Mean Absolute Error')
```



```
[38]: # Mean Absolute Percentage Error Visualization
trained_model_for_MAPE = ["KNN", "Decision Tree", "Random Forest"]
mape_scores = [knn_mape, dt_mape, rf_mape]

plt.subplot()
plt.bar(trained_models, mape_scores)
plt.ylabel("MAPE")
plt.title("Mean Absolute Percentage Error")
```

```
[38]: Text(0.5, 1.0, 'Mean Absolute Percentage Error')
```



1.18 Evaluate results using MSE, R^2 metrics

```
[39]: # Evaluation of KNN model
knn_mse = mean_squared_error(y_test, knn_predictions)
knn_r2 = r2_score(y_test, knn_predictions)

print("KNN Regression MSE:", knn_mse)
print("KNN Regression R^2:", knn_r2)
```

KNN Regression MSE: 2125543174.6260002

KNN Regression R^2 : 0.6796325091952962

```
[40]: # Evaluation of Decision Tree model
dt_mse = mean_squared_error(y_test, dt_predictions)
dt_r2 = r2_score(y_test, dt_predictions)

print("Decision Tree Regression MSE:", dt_mse)
print("Decision Tree Regression R^2:", dt_r2)
```

Decision Tree Regression MSE: 2053873667.1627307

Decision Tree Regression R^2 : 0.6904347269753497

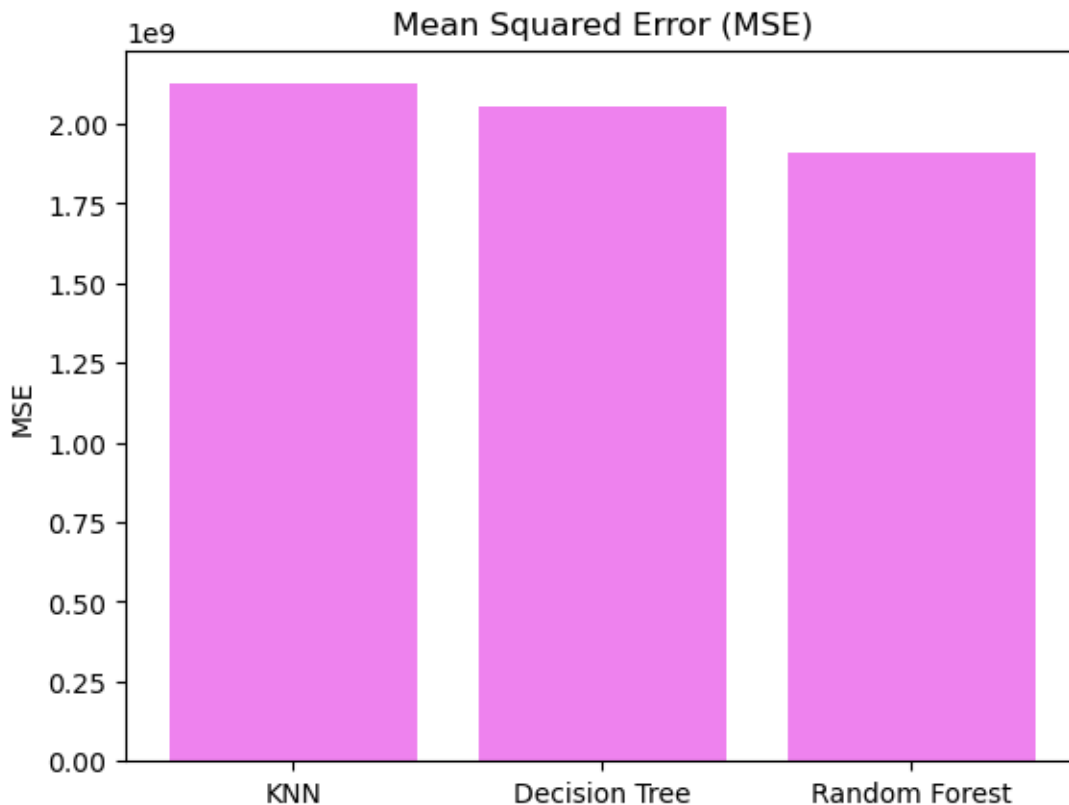

```
[41]: # Evaluation of Random Forest model
rf_mse = mean_squared_error(y_test, rf_predictions)
rf_r2 = r2_score(y_test, rf_predictions)

print("Random Forest Regression MSE:", rf_mse)
print("Random Forest Regression R^2:", rf_r2)
```

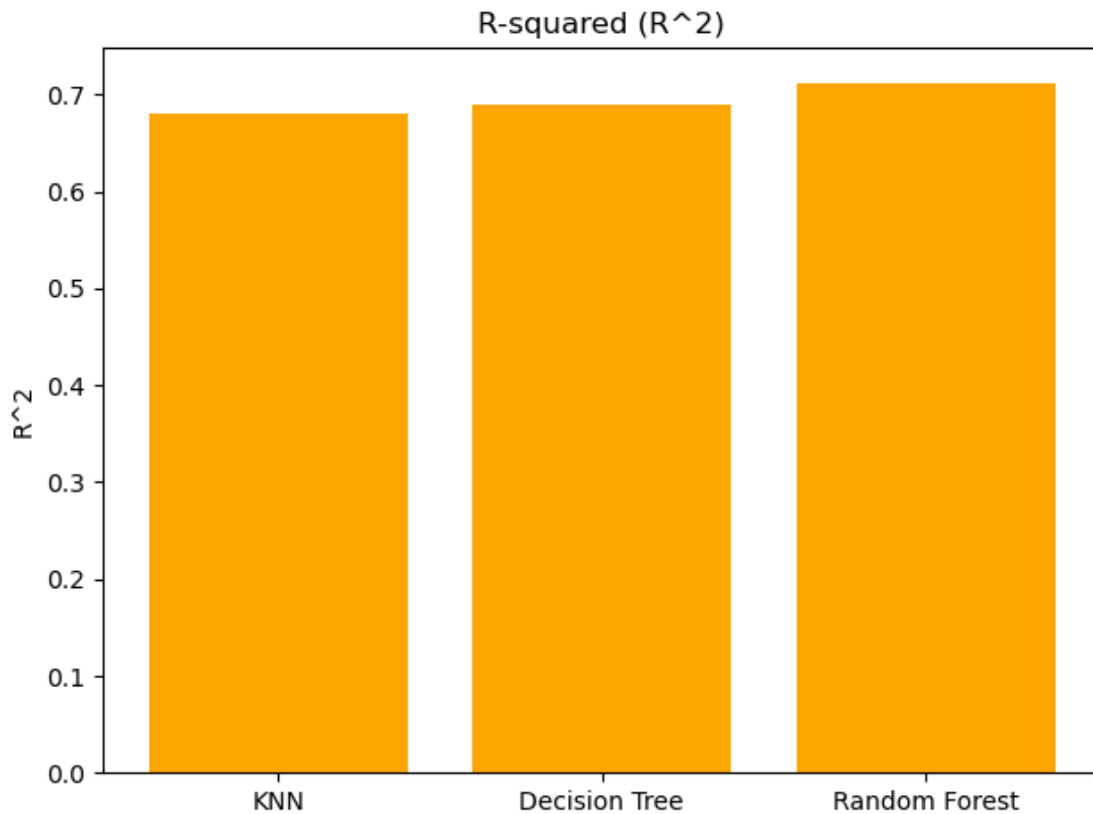
Random Forest Regression MSE: 1907485039.886032
Random Forest Regression R^2: 0.7124988081771974

```
[42]: trained_models = ['KNN', 'Decision Tree', 'Random Forest']
mse_scores = [knn_mse, dt_mse, rf_mse]
r2_scores = [knn_r2, dt_r2, rf_r2]
plt.subplot()
plt.bar(trained_models, mse_scores, color='violet')
plt.ylabel('MSE')
plt.title('Mean Squared Error (MSE)')
```

```
[42]: Text(0.5, 1.0, 'Mean Squared Error (MSE)')
```



```
[43]: plt.subplot()
plt.bar(trained_models, r2_scores, color='orange')
plt.ylabel('R^2')
plt.title('R-squared (R^2)')
plt.tight_layout()
plt.show()
```



1.19 House Price Prediction based on one-sample using Random Forest

```
[44]: # The X_test represents the testing data
# The array or list takes the rows as the input
sample_data = X_test[[3,4,40]]
predict_price = random_forest_regressor.predict(sample_data)
print("The predicted price for the sample is:", predict_price)
```

The predicted price for the sample is: [297570.81 146016.69 119605.89]

1.20 Conclusion

- Based on the above results, Random Forest model has lower rate of Mean Squared Error (Squared difference between the actual value and the predicted value) with value 1.8 compared

to Decision Tree and KNN Models.

- [1] From R-Squared evaluation metric perspective, Random Forest produced 0.7 units compared to KNN = 0.67, and Decision Tree = 0.41 which indicates Random Forest model analyzed there is a high correlation between features which explains there is small difference between the actual value and fitted value or predicted value.

Reference:

[1] <https://statisticsbyjim.com/regression/interpret-r-squared-regression/>